

Thricing Is All You Need

Identifying Risk-Neutral Densities

from an Empirical Measure of Listed Quotes

Simon A. Broda^{a,*}

*^aInstitute of Financial Services IFZ, Lucerne University of Applied Sciences, Suurstoffi
1, Rotkreuz, 6343, Switzerland*

Abstract

European option prices of a single maturity encode the risk-neutral density of the underlying at expiry. Differentiating those prices twice is ill-posed on a finite strike list. The calls are rescaled to a complementary distribution, the empirical measure of the risk-neutral jumps in strike, and that measure is fit by a smoothing spline. A Gaussian damping of its characteristic function inverts in closed form and, after the unquoted tails are filled and the kernel is thriced, supplies both the density and the call. The penalty is zero when the thriced density is already unimodal. The estimator is in-fill consistent. On four listed chains the even/odd hold-out error is smaller than each comparison method on every chain, and each fitted density has one peak, mass 1.0, and total variation 1.0. On noisy Heston quotes the mean integrated squared error is below the positive convolution approximation.

Keywords: risk-neutral density, option prices, kernel estimation, in-fill consistency, state-price density

2010 MSC: 62G07, 62G20, 62P05, 91G20

JEL: C14, C58, G13

1. Introduction

Breeden and Litzenberger (1978) identified the risk-neutral density of the underlying at a fixed expiry with the second strike derivative of the corresponding

^{*}Corresponding author.

Email address: `simon.broda@hslu.ch` (Simon A. Broda)

European call, scaled by a discount factor. The identity assumes a twice differentiable price curve. A market supplies a finite ordered list of strikes and a quote at each knot. Numerical second differentiation of that list is ill-posed: interpolants oscillate, change sign, and depend on the interpolating scheme. The literature responds with parametric mixtures and Gram–Charlier expansions, implied-volatility interpolants, local polynomials of the call, and Bondarenko (2003)’s positive convolution approximation, which forces the density to be a mixture of a positive kernel with a discrete probability fitted to prices. Section 2 reviews these approaches.

This paper estimates that density from the empirical measure of a single maturity. The calls are rescaled to a complementary distribution, which records the risk-neutral jumps in strike. A smoothing spline fits the measure: a natural cubic, with a penalty on the second derivative that is zero when the thrice density is already unimodal and positive when that derivative is noisy. A Gaussian damping of the characteristic function of the spline inverts in closed form. The same smoothing is the density and, after one integration, the call. Tails complete the unquoted wings, and thricing raises the order of the kernel. With those pieces in place, Theorem 1 gives in-fill consistency on the interior of a compact interval whenever the mesh is quasi-uniform and fills the interval, $h \rightarrow 0$ and $nh^5 \rightarrow \infty$, and the penalty is drawn from a fixed finite grid. On the exact designs below the selected penalty is zero.

We compare the fit to several extant methods. Bondarenko’s method convolves a positive kernel with a discrete probability, so the estimate is a density by construction. Aït-Sahalia and Duarte (2003) smooth a decreasing-convex projection of the calls with a local polynomial, so the density is nonnegative by design and the call cannot match quotes that violate convexity. Aït-Sahalia and Lo (1998) smooth implied volatility, and Grith et al. (2012) use a local cubic of the call and of implied volatility. On the listed chains of Section 8 the even/odd hold-out error is smaller than each of these methods on every chain, and the fitted density has one peak, mass 1.0, and total variation 1.0. The out-of-the-money root-mean-square error (RMSE) is below the convolution of Bondarenko, and on noisy Heston quotes the mean integrated squared error is below it as well.

The remainder of this manuscript is structured as follows. Section 2 reviews related work. Section 3 constructs the estimator. The characteristic function of the normalized call supplies the closed form. Section 4 records two refinements: completing the tails and thricing. Section 5 proves in-fill consistency of the estimator and states the bandwidth used below. Section 6 evaluates the convolution by a fast Fourier transform and records the speedup against the naive cell inte-

grals. Section 7 compares that estimator with those of Section 8 on exact Heston and variance-gamma quotes and on noisy Heston quotes. Section 8 presents an empirical illustration on four listed chains. Section 9 concludes.

2. Related Work

The identity of Section 3 is the continuous-strike version of a spanning argument of Ross (1976): if one could trade a put at every strike, the second derivative of the put curve would be a density, and one could synthesize any sufficiently regular payoff. Dupire (1994) obtained the analogous result in calendar time, which produces a local-volatility equation. Observed markets supply neither a continuum of strikes nor noiseless quotes.

A first family of methods specifies the density of S_T as a parametric curve and fits the implied option prices to the quotes. Ritchey (1990) and Melick and Thomas (1997) take the density to be a mixture of lognormals, so that the call price is a mixture of Black–Scholes prices; Bahra (1997) and Söderlind and Svensson (1997) use related specifications for interest-rate and foreign-exchange options. Two or three components typically suffice to generate a unimodal, skewed density, but the fit is poorly identified once more than two components are admitted, and the approximation to a Heston or jump-diffusion density need not vanish as more quotes arrive. Expansion methods perturb a baseline density—usually lognormal—by Gram–Charlier or Hermite terms (Jarrow and Rudd, 1982; Corrado and Su, 1996; Rompolis and Tzavalis, 2007; Xiu, 2014). The expansions can match a finite number of moments, but they are not guaranteed to stay nonnegative, and the truncated series can oscillate in the tails.

A second, and in applied work the most common, family interpolates the Black–Scholes implied-volatility function of strike and converts the interpolant back to prices before differentiating twice. Shimko (1993) fitted a quadratic in strike; Malz (1997) used a cubic spline in delta; Bliss and Panigirtzoglou (2002) a smoothing spline. Figlewski (2010) completes the tails with a generalized extreme-value distribution. Because implied volatility is a much flatter function of strike than the call itself, a low-order interpolant often has small bias when quotes are accurate. The drawback is that implied volatility is a nonlinear transformation that inflates price noise where the derivative of price with respect to implied volatility is small, i.e. in the tails. An interpolating spline then treats those noisy tail quotes as exact, and the second derivative of the resulting call curve oscillates and can become negative. Shape-constrained smoothers of the implied-volatility surface (Fengler, 2005, 2009; Gatheral and Jacquier, 2014;

Ludwig, 2015) address the sign problem at the cost of a more involved optimization.

A third family estimates the call, or its Black–Scholes implied volatility, by kernel regression or local polynomials and reads the density off a second strike derivative. Aït-Sahalia and Lo (1998) use a Nadaraya–Watson kernel. The state-price densities in that paper are estimated from 14,431 S&P 500 options pooled across every trading day of 1993, because a single expiry was treated as too short a list on which to differentiate twice. They emphasise that the bandwidth dominates the choice of kernel: too small a neighborhood reproduces the noise in the quotes, and too large a neighborhood removes the curvature of the smile. Aït-Sahalia and Lo (2000) used the same estimator to recover an implied risk-aversion function. Aït-Sahalia and Duarte (2003) work one day at a time, on the twenty to fifty strikes of a single expiry, and impose monotonicity and convexity of C in K . The local polynomial that the asymptotics recommend for a second derivative is a local cubic (Fan and Gijbels, 1996). On 242 days in 1999, with the 25 most active strikes, unconstrained Nadaraya–Watson, local linear, local quadratic, and local cubic estimates of the density took a negative value on 242, 130, 205, and 241 days. The shape-constrained local linear never did. Stopping the violations in the unconstrained fits required substantial oversmoothing and a large bias. The raw quotes themselves broke convexity on 24% of the high-frequency prints that year.

Grith et al. (2012) survey this literature and estimate a local cubic of the rescaled call with a quartic kernel. For a second derivative the optimal bandwidth is of order $n^{-1/9}$. On their one-month DAX chain, least-squares cross-validation selects $h = 114$, which they describe as wiggly, so the densities they plot use $h = 228$ and $h = 434$. The implied-volatility smoother they display is widened further, to $h = 372$. Smoothing the call is more sensitive to that choice than smoothing implied volatility, and the tails oscillate unless the bandwidth is increased at a cost in bias. Twice differentiating a Nadaraya–Watson or Gasser and Müller (1979) level fit is set aside as the statistically worse route to the same derivative. Jackwerth (2004) starts a kernel regression of implied volatility from Silverman’s rule, of order $n^{-1/5}$. He warns that the bandwidth which prices overfits noise from one day to the next, that a kernel does not continue across gaps in the strike list, and that a spline has to be of degree higher than three if the recovered distribution is to be smooth. Figlewski (2018) finds that filling the implied-volatility curve with a cubic spline puts spikes in the density, because the second derivative of a cubic is continuous but not differentiable, and recommends a spline of degree four or higher. Fan and Mancini (2009) combine a parametric guide with a nonparametric correction. Related shape-constrained least squares are studied by Yatchew

and Härdle (2006) and Birke and Pilz (2009).

The estimator of Section 3 is a kernel smooth of that empirical measure on a single maturity. The pilot is a smoothing spline of the complementary distribution, the kernel is Gaussian, and the bandwidth is a fixed multiple of $n^{-1/9}$. The penalty in (8) is what responds when the second derivative is noisy. Thricing raises the order of the kernel and puts negative sidelobes in the equivalent kernel for $f_{\mathbb{Q}}$. The convolution is closed form. Section 8 implements the Nadaraya–Watson implied-volatility regression of Aït-Sahalia and Lo at a bandwidth chosen there rather than in their Appendix A, the local cubics of Grith et al. at the leave-one-out width and at twice that width, and the two-step procedure of Aït-Sahalia and Duarte: the same projection of the quoted calls, followed by one Gaussian local linear at the Fan–Gijbels plug-in. The density is the scaled and shifted derivative of that local linear.

Bondarenko (2003) proposed the positive convolution approximation (PCA). The density is required to lie in the set of functions obtained by convolving a discrete probability with a positive kernel, and the discrete probability is chosen to minimize the sum of squared pricing errors. The kernel enforces smoothness and, together with positive weights, guarantees that the result is a density. In Bondarenko’s comparison, the method dominated mixtures of lognormals, Hermite expansions, and several regularization methods. Subsequent simulation evidence on affine jump-diffusions is consistent with this ranking (Lu, 2019). Section 8 implements that convolution with a Gaussian kernel on the spot. The mixing nodes are equally spaced from the first quoted strike to the last, the weights are a probability with mean F , and the bandwidth is chosen by even/odd pricing error.

Rubinstein (1994) and Jackwerth and Rubinstein (1996) recovered probabilities on a binomial tree by minimizing a distance to a prior subject to pricing constraints. Stutzer (1996) and Buchen and Kelly (1996) maximized entropy. These methods produce a valid density by construction, but the finite-support discretization and the choice of prior affect the tails. Viewing the problem as a linear inverse problem, Cont and Tankov (2004) calibrated jump-diffusion models under a relative-entropy penalty; Monnier (2014) used the singular value decomposition of a restricted put operator.

A number of papers compare methods on market data, where the true density is unknown, so that the criterion is typically pricing error or stability over time (Cooper, 1999; Jondeau and Rockinger, 2000; Coutant et al., 2001; Anagnou et al., 2005; Bliss and Panigirtzoglou, 2002; Liu et al., 2007). Studies that generate options from a known process are less common. Bondarenko (2003) finds that his method outperforms mixtures, expansions, and several regulariza-

tion methods on integrated squared error. Lai (2014) finds that kernel regression outperforms spline interpolation and neural nets on pricing error. Santos and Guerra (2015) compare mixtures, Edgeworth expansions, hypergeometric functions, and a Bliss and Panigirtzoglou smoothing spline against Heston and other targets. Lu (2019) compares methods on affine jump-diffusions and again finds Bondarenko's method accurate. The typical ranking is that implied-volatility interpolants price well when quotes are clean, that mixtures are unstable once more than two components are admitted, and that Bondarenko's method is the most accurate among methods that do not use a parametric density of S_T , when the true density is known. Section 7 is in this vein: the targets are Heston and variance-gamma, and the true density is known. Section 8 scores listed index chains, where the criterion is pricing error, against Bondarenko, Aït-Sahalia and Duarte, Aït-Sahalia and Lo, and Grith et al..

3. Construction of the Estimator

3.1. Construction

Assume a risk-neutral measure $\mathbb{Q}$ under which S_T has a density $f_{\mathbb{Q}}$ on $(0, \infty)$ and $\mathbb{E}^{\mathbb{Q}}[S_T] < \infty$. Let r be the continuously compounded interest rate and q the continuous dividend yield. Write $F = S_0 e^{(r-q)T}$ for the forward. Under $\mathbb{Q}$,

$$\begin{aligned} S_0 e^{-qT} &= e^{-rT} \mathbb{E}^{\mathbb{Q}}[S_T] \\ &= e^{-rT} \int_0^{\infty} x f_{\mathbb{Q}}(x) dx, \end{aligned} \tag{1}$$

hence $\mathbb{E}^{\mathbb{Q}}[S_T] = F$. Write C for the corresponding call curve:

$$\begin{aligned} C(K) &= e^{-rT} \mathbb{E}^{\mathbb{Q}}[(S_T - K)^+] \\ &= e^{-rT} \int_K^{\infty} (x - K) f_{\mathbb{Q}}(x) dx. \end{aligned}$$

Now define the normalized option price

$$\tilde{C}(K) \equiv \frac{C(K)}{S_0 e^{-qT}} = \frac{\int_K^{\infty} (x - K) f_{\mathbb{Q}}(x) dx}{F}. \tag{2}$$

The normalized call price is a decreasing function of K : it is equal to 1 if $K = 0$ and equal to 0 if $K = \infty$. Thus the normalized call price, as a function

of K , is a complementary cdf, say $\tilde{C}(K) = 1 - P(K)$. The characteristic function (hereafter c.f.) of the law of P is the Fourier–Stieltjes transform

$$\varphi_P(s) = \int_{(0,\infty)} e^{isK} dP(K). \quad (3)$$

For a cdf on $[0, \infty)$ and $s \neq 0$, (3) is

$$\varphi_P(s) = 1 + is \int_0^\infty e^{isK} \tilde{C}(K) dK.$$

Substitute (2). Fubini yields

$$\int_0^\infty e^{isK} \tilde{C}(K) dK = \frac{1}{F} \int_0^\infty f_{\mathbb{Q}}(x) \left[\int_0^x e^{isK} (x - K) dK \right] dx.$$

The inner integral is $-x/(is) + (e^{isx} - 1)/(is)^2$. Hence, for $s \neq 0$,

$$\begin{aligned} \varphi_P(s) &= 1 + \frac{is}{F} \int_0^\infty f_{\mathbb{Q}}(x) \left[-\frac{x}{is} + \frac{e^{isx} - 1}{(is)^2} \right] dx \\ &= 1 - F^{-1} \mathbb{E}^{\mathbb{Q}}[S_T] + \frac{\varphi_{\mathbb{Q}}(s) - 1}{is F} \\ &= \frac{\varphi_{\mathbb{Q}}(s) - 1}{is F}, \end{aligned} \quad (4)$$

where $\varphi_{\mathbb{Q}}(s) = \mathbb{E}^{\mathbb{Q}}[e^{isS_T}]$ and $\mathbb{E}^{\mathbb{Q}}[S_T] = F$ by (1). Solving for $\varphi_{\mathbb{Q}}$, one obtains

$$\varphi_{\mathbb{Q}}(s) = 1 + is F \varphi_P(s). \quad (5)$$

Read from the right, (4) is a pricing formula: a model that supplies $\varphi_{\mathbb{Q}}$ yields φ_P , hence the complementary cdf $\tilde{C} = 1 - P$ and the call (Lewis, 2000). This paper inverts the identity: φ_P is constructed from the quotes, and $\varphi_{\mathbb{Q}}$ is recovered.

The c.f. $\varphi_{\mathbb{Q}}$ is that of S_T . Fourier inversion yields

$$f_{\mathbb{Q}}(K) = \frac{1}{2\pi} \int_{-\infty}^\infty e^{-isK} \varphi_{\mathbb{Q}}(s) ds. \quad (6)$$

Substitute (5) into (6):

$$f_{\mathbb{Q}}(K) = \frac{1}{2\pi} \int_{-\infty}^\infty e^{-isK} ds + \frac{F}{2\pi} \int_{-\infty}^\infty e^{-isK} (is) \varphi_P(s) ds.$$

The first integral is a Dirac mass at $K = 0$. For $K > 0$,

$$f_{\mathbb{Q}}(K) = \frac{F}{2\pi} \int_{-\infty}^{\infty} e^{-isK} (is) \varphi_P(s) ds. \quad (7)$$

Observed calls at strikes $0 < K_1 < \dots < K_m$ supply values $\tilde{C}_i = C_i/(S_0 e^{-qT})$ and hence $P_i = 1 - \tilde{C}_i$. The quotes are noisy, and thricing amplifies the second derivative, so the estimator does not pass a cubic through every quote. It fits a smoothing spline, the natural cubic f that minimises

$$\sum_i (P(K_i) - f(K_i))^2 + c F^3 \int (f''(K))^2 dK. \quad (8)$$

The penalty is zero when the thriced density is already unimodal and positive when that derivative is noisy. The factor F^3 makes c dimensionless. Section 8.2 chooses c from a fixed finite grid. On the exact Heston and variance-gamma quotes below, that rule selects zero. The minimiser, written P^{cs} , is a cubic spline, so the integrals below are closed form. The c.f. of this absolutely continuous measure is $\int e^{isK} P^{cs'}(K) dK$. Replace φ_P by the product of that integral with the Fourier transform of a Gaussian kernel of width h ,

$$\hat{\varphi}_P(s) = \left(\int_{(0,\infty)} e^{isK} P^{cs'}(K) dK \right) e^{-h^2 s^2 / 2}. \quad (9)$$

The Gaussian density and its distribution function are

$$\kappa_h(u) = \frac{1}{h \sqrt{2\pi}} \exp\left(-\frac{u^2}{2h^2}\right), \quad \mathcal{K}_h(u) = \int_{-\infty}^u \kappa_h.$$

The factor $e^{-h^2 s^2 / 2}$ is convolution against κ_h , and multiplication by is is differentiation up to sign. Substituting (9) into (7) therefore yields

$$\hat{f}_{\mathbb{Q}}(K) = -F (P^{cs'} * \kappa_h)'(K), \quad \hat{P}(K) = (P^{cs'} * \mathcal{K}_h)(K). \quad (10)$$

On each cell the integrand is a quadratic times a Gaussian. The integrals of 1, u , and u^2 against ϕ and Φ are elementary, so (10) is closed form. Section 6 evaluates the same convolution from (9) with a fast Fourier transform. Calls are $\hat{C} = S_0 e^{-qT} (1 - \hat{P})$. Section 4.1 adds the endpoints that complete the distribution, and those knots enter the sum in (8).

4. Two Refinements

4.1. Completing the Tails

The spline of Section 3.1 stops at the first and last quotes. On $(0, K_1)$ it does not yet specify P , and if $C_m > 0$ the residual mass $1 - P(K_m)$ is unplaced. Both ends are wrong as a cdf on $(0, \infty)$.

On $(0, K_1]$ a deep in-the-money call is intrinsic, $C(K) = e^{-rT}(F - K)$, hence $\tilde{C}(K) = 1 - K/F$ and $P(K) = K/F$. That is a straight line through the origin. The estimator matches the observed $P(K_1)$ rather than K_1/F ,

$$P(K) = P(K_1) \frac{K}{K_1}, \quad 0 < K \leq K_1. \quad (11)$$

When the first quote is at intrinsic, $P(K_1) = K_1/F$ and (11) is exactly K/F . The fit includes the knot $P(0) = 0$, which is the same line at the boundary.

On the right, $\tilde{C}(\infty) = 0$ so $P(\infty) = 1$. If $C_m > 0$, a European call still has to fall to zero. The residual call is continued linearly with slope in $[-e^{-rT}, 0)$ until $C = 0$ at an intercept K_{end} . The slope is the last strictly negative step of the quoted call, clipped to that interval. A flat last step, two quotes on the premium floor, is not a decay: using it would send K_{end} to the cap $8F$. The fit includes the knot $P(K_{\text{end}}) = 1$. The bandwidth still counts a notional knot spacing along that line, eight per last quoted gap,

$$n_R = \text{round}(8(K_{\text{end}} - K_m)/(K_m - K_{m-1})), \quad (12)$$

are equally spaced in $(K_m, K_{\text{end}}]$ and converted to $P = 1 - \tilde{C}$. If $C_m = 0$, $P(K_m) = 1$ already and this wing is empty.

Only the left completion changes the estimator on the Heston and variance-gamma chains below: the first quoted call is at intrinsic, and the last quoted call is already zero (dense) or of order 10^{-5} (sparse).

4.2. Thricing

The Gaussian kernel κ_h has second moment h^2 , so the leading bias of (10) is $O(h^2)$. Schucany and Sommers (1977) call the first correction twicing. It replaces $e^{-h^2 s^2/2}$ by $2e^{-h^2 s^2/2} - e^{-h^2 s^2}$, the transform of $2\kappa_h - \kappa_{h\sqrt{2}}$. The second moment vanishes, and the bias falls from $O(h^2)$ to $O(h^4)$. We take that construction one step further. A third bandwidth cancels the new leading term, and we call the result thricing. The transform in (9) is replaced by

$$\frac{8}{3}e^{-h^2 s^2/2} - 2e^{-h^2 s^2} + \frac{1}{3}e^{-2h^2 s^2}, \quad (13)$$

which is the Fourier transform of $\frac{8}{3}\kappa_h - 2\kappa_{h\sqrt{2}} + \frac{1}{3}\kappa_{2h}$. The moments of order 2 and 4 vanish, so the bias is $O(h^6)$ when $f_{\mathbb{Q}}$ has a bounded sixth derivative. In closed form,

$$\hat{f}_{\mathbb{Q}}^{\text{th}}(K) = \frac{8}{3}\hat{f}_{\mathbb{Q},h}(K) - 2\hat{f}_{\mathbb{Q},h\sqrt{2}}(K) + \frac{1}{3}\hat{f}_{\mathbb{Q},2h}(K), \quad (14)$$

and the same combination is applied to $\hat{P}$ in (10). The equivalent kernel takes negative values, so $\hat{f}_{\mathbb{Q}}^{\text{th}}$ is not a density. The call is $S_0 e^{-qT}(1 - \hat{P}^{\text{th}})$, with $\hat{P}^{\text{th}}$ clipped to $[0, 1]$ and both premiums floored at zero.

The correction is large because a call is an integral of $f_{\mathbb{Q}}$. Thus P is smoother than the density by two derivatives, and the bias expansion is the expansion of that integral against a Gaussian. A bandwidth of order $n^{-1/9}$ balances an $O(h^2)$ bias against a variance of order $1/(nh^5)$, so on exact quotes that bias is the leading error. Twicing cancels it, and thricing cancels the $O(h^4)$ term that replaces it. The negative sidelobes have width proportional to h . On the quoted strikes alone the natural spline sets $P'' = 0$ at the extreme quotes, inside the window the density is read from. The knots of Section 4.1 move that boundary to 0 and to K_{end} , so the sidelobes integrate the completed tails. On a truncated strip the same kernel amplifies the missing mass, which is why thricing is applied only after that completion.

5. In-Fill Consistency

Write P for the complementary-call cdf of Section 3, so that $f_{\mathbb{Q}} = -F P''$ on $(0, \infty)$. The estimator is the thriced smoothing spline (14).

Theorem 1. *Let $\mathbb{E}^{\mathbb{Q}}[S_T] < \infty$, and let $f_{\mathbb{Q}}$ be four times continuously differentiable on $(0, \infty)$. Let C be the corresponding call curve. Fix a compact interval $[A, B] \subset (0, \infty)$. Consider a sequence of quote grids $0 < K_1^{(m)} < \dots < K_m^{(m)}$ with mesh $\delta_m = \max_i(K_i^{(m)} - K_{i-1}^{(m)}) \rightarrow 0$, $K_1^{(m)} \leq A$ and $K_m^{(m)} \geq B$, and with the mesh on $[A, B]$ quasi-uniform: $n_m \delta_m$ bounded above and below, where n_m is the number of strikes in $[A, B]$. Observe*

$$P_i^{(m)} = P(K_i^{(m)}) + \eta_i^{(m)},$$

with $\eta_i^{(m)}$ independent, $\mathbb{E}[\eta_i^{(m)}] = 0$, and $\mathbb{E}[(\eta_i^{(m)})^2] \leq \tau^2 < \infty$. Let the penalty c in (8) be any random variable with values in a fixed finite set $C \subset [0, \infty)$, and let F stay fixed. Form the smoothing spline of Section 3.1, the tails of Section 4.1, and (14). Let $h_m \rightarrow 0$ and $n_m h_m^5 \rightarrow \infty$. Then $\hat{f}_{\mathbb{Q}}^{\text{th}}$ converges in probability to $f_{\mathbb{Q}}$, uniformly on every compact subset of (A, B) .

Proof. The tail knots lie outside $[A, B]$. On a compact several bandwidths inside (A, B) their contribution is of order $h^{-2} \exp(-\varepsilon^2/2h^2)$, whether or not the intercept K_{end} depends on the last noisy slope, because that intercept stays inside a fixed multiple of F .

For each fixed penalty the smoothing spline is linear in the knots, and F is fixed, so $\lambda = cF^3$ has finitely many values. A maximum of finitely many $o_p(1)$ terms is $o_p(1)$. Fix $c \in C$. The fitted spline is the spline of the true knots plus the spline S_η of $(\eta_i^{(m)})$.

Let P_0 be the spline of the true knots at penalty zero and P_λ the spline of the same knots at this λ . On a quasi-uniform mesh the natural cubic through a C^4 function converges to the function, and its first derivative to the first derivative, uniformly on interior compacts, so $P'_0 \rightarrow P'$. The thrice kernel integrates to one and is even. Convolution against it, followed by one derivative, sends P' to P'' as $h \rightarrow 0$, and therefore sends P'_0 to P'' , uniformly on those compacts. When f_Q is six times continuously differentiable the moments of order 2 and 4 vanish and that bias is $O(h^6)$.

The variational inequality gives $\sum_i (P(K_i) - P_\lambda(K_i))^2 \leq \lambda \int (P''_0)^2$. The integral stays bounded because $P''_0 \rightarrow P''$ on the interior of the grid. The difference $P_0 - P_\lambda$ is the spline through these residuals. The condition $n_m h_m^5 \rightarrow \infty$ forces $h_m/\delta_m \rightarrow \infty$ on a quasi-uniform mesh. The contribution of the residuals to the differentiated convolution is then within $o(1)$ of

$$\sum_i r_i \delta_i h_m^{-3} L((K - K_i)/h_m),$$

and Cauchy–Schwarz bounds the sum by a constant times $(\lambda/(n_m h_m^5))^{1/2} \rightarrow 0$.

On the noise, the Demmler–Reinsch basis multiplies the coefficient of the j th function by $(1 + \lambda \rho_j)^{-1}$ with $\rho_j \geq 0$, so every linear functional of S_η has variance at most its variance at penalty zero. At penalty zero, $(S'_\eta * \kappa_{h_m})'$ is within $o_p(1)$ of

$$\sum_i \eta_i^{(m)} \delta_i h_m^{-3} L((K - K_i)/h_m),$$

where L is a fixed integrable combination of second derivatives of the thrice kernel and $\int L = 0$. The sum has mean zero and conditional variance of order $\tau^2/(n_m h_m^5)$, uniformly on the compact. Chebyshev's inequality gives pointwise convergence in probability, and a Lipschitz constant of the same order upgrades it to uniform convergence.

Multiplying by $-F$ yields $\hat{f}_Q^{\text{th}} \rightarrow f_Q$. □

Remark 1. *The theorem allows additive noise in P and a penalty drawn from a fixed finite grid. Exact quotes are the case $\eta_i = 0$. The limit is uniform on compact subsets of (A, B) . The equivalent kernel of (14) takes negative values, so the estimator is not necessarily a density. The empirical section fits the spline to the decreasing-convex projection of the quoted calls. When the true call lies in the interior of that cone, the projection changes the knots by an amount of the same stochastic order as η_i , and the bandwidth condition of Theorem 1 still applies.*

The bandwidth h in (14) is of order $n^{-1/9}$. That rate balances an $O(h^2)$ bias against a variance of order $1/(nh^5)$, and any fixed positive multiple meets Theorem 1. The kernel-density exponent $n^{-1/5}$ leaves nh^5 bounded, so it is not used. Write $\hat{\sigma}_{\text{ATM}}$ for the at-the-money Black–Scholes implied volatility, the median of contracts with $|K - F| \leq 0.03F$ (or the median of all finite implied volatilities if that window is empty), and take

$$h = 0.38 F \hat{\sigma}_{\text{ATM}} \sqrt{T} n^{-1/9}, \quad (15)$$

with n the number of quoted strikes. This 0.38 is the bandwidth constant. The penalty in (8) is a different constant, chosen later. This is the bandwidth used below.

6. Computational concerns

The product in (9) is the convolution (10). The factor $e^{-h^2 s^2/2}$ damps the characteristic function of the spline slope, multiplication by is differentiates that slope, and division by is integrates it. Thricing replaces the damping by (13). Each step multiplies one transform, so a single discrete Fourier transform evaluates the plain convolution and the thrice combination together.

The completed spline is sampled on a uniform grid and held at its endpoint values outside the knots. The slope in (10) is the derivative of that extension, so the transform of the slope is is times the transform of the samples. The grid step is one sixty-fourth of the smaller of the median knot spacing and the bandwidth. The grid runs ten widths of the widest kernel past every strike at which the density or the call is read, which keeps the periodic transform from wrapping into that window. The level in (10) is the damped transform with the left endpoint removed. Every quoted strike is interpolated from that grid, so one transform of the chosen spline prices the chain.

The same convolution is the sum, over the cells of the spline, of the elementary integrals in Section 3.1. On the exact, noisy, and listed designs below, the Fourier

evaluation reproduces the printed integrated squared errors, the selected penalties, and the listed out-of-the-money errors. Table 1 records the wall-clock time of one fit at (15), the search over the penalty grid included. Against that naive sum over cells, the Fourier evaluation is between 13.8 and 95.2 times as fast.

Table 1: Wall-clock seconds for one fit of the estimator. Naive is the sum of the closed-form cell integrals. FFT is the discrete Fourier evaluation of (9) and (13). The speedup is the ratio of those times.

Design	Naive	FFT	Speedup
Heston, dense	4.317	0.049	88.6
Heston, sparse	0.244	0.013	18.4
Variance gamma, dense	4.182	0.044	95.2
Variance gamma, sparse	0.238	0.017	13.8
SPX 18 Dec 2026	8.483	0.122	69.4
SPX 19 Mar 2027	3.654	0.072	51.1
NDX 18 Dec 2026	5.488	0.071	77.8
RUT 18 Dec 2026	1.498	0.058	25.8

7. Numerical Illustration

The data-generating process is the stochastic-volatility model of Heston (1993) with the equity parameters of Bakshi et al. (1997): $S_0 = 100$, $r = 0.05$, $q = 0$, $T = 1/4$, $\kappa = 1.15$, $\theta = 0.04$, $\sigma = 0.39$, $\rho = -0.64$, $v_0 = 0.04$. The c.f. of $\log S_T$ uses the little-trap branch of Albrecher et al. (2007). Call prices are the damped integrand of Carr and Madan (1999) with damping $\alpha = 1.5$, evaluated by the trapezoidal rule on 4096 frequencies in $(0, 250]$; puts follow by parity. The true density is Fourier inversion of the same c.f. on the same frequency grid,

$$f_{\mathbb{Q}}(K) = \frac{1}{K\pi} \int_0^\infty \operatorname{Re}[e^{-iu \log K} \varphi(u)] du,$$

where $\varphi(u) = \mathbb{E}^{\mathbb{Q}}[e^{iu \log S_T}]$. Two chains are used: a dense grid of $m = 256$ equally spaced strikes on $[30, 220]$ (spacing $\delta = 0.745$) and a sparse grid of $m = 16$ equally spaced strikes on $[70, 140]$ (spacing $\delta = 4.67$). Integrated squared error (ISE) is $\int_{60}^{150} (\hat{f}_{\mathbb{Q}} - f_{\mathbb{Q}})^2 dK$, evaluated by the trapezoidal rule on 401 equally spaced points. Table 2 reports ISE for four nested estimators at the oracle h^* and at (15). The oracle minimizes the integral on a geometric grid in h .

Table 2: Integrated squared error on exact Heston quotes for the cubic-spline estimator. The last pair is (15). Quoted spline uses only the observed strikes. Tails adds the wings of Section 4.1. Tails+Twice is twicing (Schucany and Sommers, 1977). Tails+Thrice applies (14).

Chain	Estimator	h^*	Oracle ISE	$h_{1/9}$	$h_{1/9}$ ISE
Dense	Quoted spline	0.15	9.76×10^{-10}	2.00	1.70×10^{-5}
Dense	Tails	0.15	9.76×10^{-10}	2.00	1.70×10^{-5}
Dense	Tails+Twice	0.39	2.68×10^{-12}	2.00	2.50×10^{-7}
Dense	Tails+Thrice	0.63	1.16×10^{-12}	2.00	2.50×10^{-8}
Sparse	Quoted spline	16.35	1.80×10^{-2}	2.66	1.05×10^{-1}
Sparse	Tails	0.93	3.40×10^{-6}	2.66	5.08×10^{-5}
Sparse	Tails+Twice	1.91	1.68×10^{-6}	2.66	3.13×10^{-6}
Sparse	Tails+Thrice	2.20	1.43×10^{-6}	2.66	1.60×10^{-6}




Figure 1: Heston density of S_T on $[60, 150]$. Quoted spline, tails, tails plus twice, and tails plus thrice from Table 2, each at (15). Sparse quoted strikes are marked on the true density. Curves are truncated at zero, and the vertical scale is that of the true density.

On the dense grid the first strike lies left of the window and the last call is already zero, so the wings do not move the error. At (15), twicing lowers the tails error and thricing lowers it further. That bandwidth is wider than the oracle, and the error is larger. On the sparse grid the quoted spline without tails misses the window. The left line (11) restores it, and each of the two higher-order kernels then reduces the error. Thricing at (15) sits close to its oracle. Figure 1 plots the quoted spline, tails, tails plus twice, and tails plus thrice at (15).

Table 3 repeats the design for the variance-gamma process of Madan et al. (1998). The process parameters are those of Carr and Madan (1999): $\sigma = 0.12$,

$\nu = 0.2$, $\theta = -0.14$. Spot, rate, dividend yield, and maturity are as above, as are the two strike grids and the ISE window. Variance gamma is a pure-jump Lévy process. The true density of $\log S_T$ is a modified Bessel function (Madan et al., 1998); the density of S_T is that formula times $1/K$. Calls are the discounted payoff of that density, evaluated by the trapezoidal rule.

Table 3: Integrated squared error on exact variance-gamma quotes for the cubic-spline estimator. Columns as in Table 2.

Chain	Estimator	h^*	Oracle ISE	$h_{1/9}$	$h_{1/9}$ ISE
Dense	Quoted spline	0.17	5.12×10^{-7}	1.28	1.55×10^{-4}
Dense	Tails	0.17	5.12×10^{-7}	1.28	1.55×10^{-4}
Dense	Tails+Twice	0.33	3.93×10^{-7}	1.28	3.22×10^{-5}
Dense	Tails+Thrice	0.46	4.04×10^{-7}	1.28	1.95×10^{-5}
Sparse	Quoted spline	14.17	3.27×10^{-2}	1.65	1.71×10^{-1}
Sparse	Tails	0.93	5.78×10^{-4}	1.65	6.92×10^{-4}
Sparse	Tails+Twice	1.91	5.12×10^{-4}	1.65	5.14×10^{-4}
Sparse	Tails+Thrice	2.20	4.87×10^{-4}	1.65	5.18×10^{-4}




Figure 2: Variance-gamma density of S_T on $[60, 150]$. Quoted spline, tails, tails plus twice, and tails plus thrice from Table 3, each at (15). Sparse quoted strikes are marked on the true density. Curves are truncated at zero, and the vertical scale is that of the true density.

On the dense chain the wings again do not move the error. At (15), twicing lowers it and thricing lowers it further. That bandwidth is wider than the oracle, and the error is larger. The density of S_T is more peaked than under Heston, so the errors are larger than on that design. On the sparse chain the quoted spline

without tails misses the window. Tails restore it. Twicing lowers that error, and thricing at (15) remains near its oracle. Figure 2 plots the quoted spline, tails, tails plus twice, and tails plus thrice at (15).

Table 4 scores the estimator of Section 3.1 against the methods of Section 8 on the same four chains. Each fit starts from the decreasing-convex projection. On these exact quotes the penalty is zero, so the row is the thriced spline at (15).

Table 4: Integrated squared error on $[60, 150]$ for the estimator of Section 3.1 and the methods of Section 8. Ours is thriced at (15); on these exact quotes the penalty is zero. Each other row uses that estimator’s bandwidth from Section 8. Grith et al. is shown at the leave-one-out width and at twice that width.

Estimator	Heston dense	Heston sparse	VG dense	VG sparse
Ours	2.50×10^{-8}	1.60×10^{-6}	1.95×10^{-5}	5.18×10^{-4}
PCA	8.55×10^{-8}	7.16×10^{-3}	2.05×10^{-4}	2.55×10^{-2}
Grith et al., IV, 2×	9.27×10^{-8}	2.81×10^{-5}	1.15×10^{-4}	1.67×10^{-3}
Grith et al., IV, CV	6.24×10^{-9}	4.79×10^{-6}	2.16×10^{-5}	7.40×10^{-4}
Aït-Sahalia and Duarte	6.13×10^{-4}	2.47×10^{-3}	6.47×10^{-3}	1.95×10^{-3}
Grith et al., call, 2×	1.14×10^{-5}	5.22×10^{-3}	4.45×10^{-4}	1.89×10^{-2}
Grith et al., call, CV	7.52×10^{-7}	6.81×10^{-4}	5.39×10^{-5}	6.65×10^{-3}
Aït-Sahalia and Lo	2.20×10^{-8}	4.29×10^{-6}	3.22×10^{-5}	4.94×10^{-4}

The estimator is the smallest on sparse Heston and on dense variance gamma. It is also smaller than the positive convolution, the plug-in local linear, both call cubics, and the doubled implied-volatility cubic on every chain. On dense Heston the cross-validated implied-volatility cubic and the Nadaraya–Watson smile of Aït-Sahalia and Lo are smaller. On sparse variance gamma that smile is smaller.

The same designs are repeated with noise in the quotes. Black–Scholes implied volatility receives independent $N(0, 0.01^2)$ noise. Values outside $[0.02, 2.5]$ are clipped, and the call is rebuilt. Each replication starts from the projection of those calls, and the penalty is chosen on that draw by the rule of Section 8.2. Table 5 is that comparison.

Table 5: Mean integrated squared error on $[60, 150]$ for noisy Heston quotes. Implied volatility is perturbed by $N(0, 0.01^2)$ noise. Thirty replications. Ours is the smoothing spline of Section 3.1, at the bandwidth (15), with the penalty chosen on each draw by the rule of Section 8.2. Each other row uses that estimator’s bandwidth from Section 8.

Estimator	Dense	Sparse
Ours	2.22×10^{-4}	5.40×10^{-4}
PCA	4.36×10^{-4}	7.20×10^{-3}
Grith et al., IV, $2\times$	2.56×10^{-3}	8.61×10^{-5}
Grith et al., IV, CV	1.56×10^{-2}	2.74×10^{-4}
Aït-Sahalia and Duarte	5.95×10^{-4}	6.33×10^{-3}
Grith et al., call, $2\times$	2.17×10^{-3}	5.81×10^{-3}
Grith et al., call, CV	1.54×10^{-2}	1.15×10^{-3}
Aït-Sahalia and Lo	2.15×10^{-4}	4.30×10^{-4}

On the dense chain the Nadaraya–Watson smile is smaller, and the smoothing spline is smaller than the positive convolution and the local polynomials. On the sparse chain the implied-volatility cubics and that smile are smaller.

8. Empirical Applications

8.1. Data and Methodology

The listed comparison uses European, AM-settled index options from Cboe’s delayed-quotes feed (`cdn.cboe.com`). Each file is dated 23 September 2026. The S&P 500 (SPX) file is 03:56 Eastern (30,470 listed contracts across 58 expiries); the Nasdaq-100 (NDX) file is 03:44 Eastern (16,492 contracts across 46 expiries); the Russell 2000 (RUT) file is 03:44 Eastern (10,790 contracts across 30 expiries). Only the standard monthly roots SPX, NDX, and RUT are kept, not weekly SPXW, NDXW, or RUTW. Time to expiry is $T = d/365.25$ with d the calendar day count from the snapshot date. The discount rate is $r = 0.04$.

A quote is retained if the bid is positive, the ask is at least the bid, the mid is at least 0.25, and the spread is at most $\max(15, 0.002S_0, 0.5 \times \text{mid})$. The 15-point term is the tolerance on the S&P 500. On the Nasdaq-100 a fixed 15-point cap deletes puts whose spreads are a dollar or two wider and leaves a gap of several thousand points in the strike list. Among remaining put–call pairs, the forward is the median of $K + e^{rT}(C - P)$ over strikes with $|K - K_{\text{atm}}| \leq 50$ on SPX, where K_{atm} is the pair closest to the cash index; on NDX and RUT the band is 0.65% of the spot. Dividend yield is $q = r - \log(F/S_0)/T$. The working slice is out

of the money (OTM): the put mid if $K \leq F$, the call mid if $K \geq F$, the missing side filled by parity. Puts with $K < 0.7F$ and mid below 0.50 are then dropped; out-of-the-money calls with mid at least 0.25 are kept. Table 6 records the filter counts. In (10) and (14), $\tilde{C} = C/(S_0 e^{-qT})$ and $F = S_0 e^{(r-q)T}$.

The December 2026 S&P 500 AM quarterly ($T = 0.235$, $S_0 = 7764.64$, $F = 7836$, $q = 0.11\%$) appears in Figure 3. After the screens it has $n = 444$ strikes on $[0.36F, 1.30F]$, of which 75% are puts. Near the money the strike mesh is 5 (0.06% of F); it is 5–25 on the put wing and up to 200 beyond $1.2F$. Median out-of-the-money mids are 39 for puts and 62 for calls. Implied volatility $\hat{\sigma}_{\text{ATM}} = 0.134$ is the median Black–Scholes implied volatility of contracts with $|K - F| \leq 0.03F$. The linear call tail of Section 4.1 is appended ($C_m = 0.25$ by the mid floor of the filter, $n_R = 80$ from (12)). The bandwidth (15) uses the 444 quoted strikes and equals $h = 98.6$. The March 2027 AM quarterly ($T = 0.485$, $F = 7918$, $q = -0.04\%$, $\hat{\sigma}_{\text{ATM}} = 0.144$) filters to $n = 221$ strikes on $[0.20F, 1.49F]$, with a 25-point mesh at the money; (12) gives $n_R = 41$, and (15) gives $h = 165.9$.

The Nasdaq-100 December 2026 AM quarterly ($T = 0.235$, $S_0 = 30732$, $F = 31025$, $q = -0.03\%$, $\hat{\sigma}_{\text{ATM}} = 0.201$) is built with the same screens. It filters to $n = 312$ strikes on $[0.29F, 1.39F]$, 83% puts; the at-the-money mesh is 25, and gaps of 1000 appear beyond $1.2F$. The bandwidth (15) is $h = 606.0$ ($n = 312$; the right wing has $n_R = 72$).

The Russell 2000 December 2026 AM quarterly ($T = 0.235$, $S_0 = 2889.92$, $F = 2912$, $q = 0.80\%$, $\hat{\sigma}_{\text{ATM}} = 0.183$) filters to $n = 99$ strikes on $[0.45F, 1.32F]$, 59% puts; the at-the-money mesh is 10. The last two calls are both 0.25, so the wing slope is the previous negative step. Equation (12) gives $n_R = 40$. The bandwidth (15) is $h = 59.0$.

The quoted calls are replaced by the nearest vector in Euclidean distance that is decreasing and convex in strike. Consecutive slopes lie in $[-e^{-rT}, 0]$, and each coordinate is at least the intrinsic value $e^{-rT} \max(F - K_i, 0)$. This is the unpenalized projection of Yatchew and Härdle and also the first step of Aït-Sahalia and Duarte. Puts of the projected call follow by parity, and both premiums are floored at zero.

Out-of-the-money RMSE uses the quoted mid of the put for $K \leq F$ and that of the call for $K \geq F$. Hold-out is the average of two even/odd splits (train on even strikes and score odd, then the reverse), with h recomputed from the training n and the projection recomputed from the training strikes only.

Table 6: Construction of the working slices. “At expiry” counts listed puts and calls on that root and maturity with a positive bid and ask at least the bid. “Screens” applies the mid and spread filters. “Working OTM” is the slice used by the estimators.

Slice	At expiry	Screens	Unique K	Working OTM	Puts
SPX 18 Dec 2026	927	922	473	444	335
SPX 19 Mar 2027	453	451	228	221	165
NDX 18 Dec 2026	631	631	319	312	259
RUT 18 Dec 2026	218	218	115	99	58

8.2. Estimators

Ours. Start from the projected calls of Section 8.1. Set $P_i = 1 - C_i/(S_0 e^{-qT})$. Fit (8) with the tail knots of Section 4.1, and invert at (15) with thricing. The constant c is chosen from zero and sixteen values spaced geometrically from 10^{-7} to 3×10^{-4} . Let c_0 be the smallest candidate at which $\hat{f}_{\mathbb{Q}}^{\text{th}}$ has one peak on $[0.55F, 1.40F]$ and total variation at most 1.02. A curve that already meets that test keeps c_0 . Above c_0 , the log root-mean-square gap between the fitted calls and the input quotes is plotted against the log root-mean-square gap between $\hat{f}_{\mathbb{Q}}^{\text{th}}$ and the plain Gaussian convolution, on that same window. If the discrete curvature of the curve is at least $1/2$, the penalty is the corner. A flatter bend keeps c_0 . Prices are floored at zero, with puts by parity.

Ait-Sahalia and Duarte. Start from the projected calls of Section 8.1. The price is a Gaussian local linear regression of those values. The bandwidth is the plug-in in their Section 3.5, taken from Fan and Gijbels (1996): a global polynomial of order four supplies the pilot curvature, the weight equals one within 1.5 standard deviations of the mean strike, the Gaussian constant is 0.776, and the rate is $n^{-1/5}$. The widths are $h = 140.9, 256.1, 715.3$, and 68.8 . The density is e^{rT} times the strike derivative of the local slope, scaled so that the mass on $[\max(50, 0.2F), 2.4F]$ equals one and shifted so that the mean equals the forward. Their May 1999 figure uses the bandwidth that minimised integrated squared error in a Monte Carlo sample of 25 strikes. That optimum is not stated as a formula, and the listed row uses the plug-in. The curves have 2, 1, 4, and 4 peaks. Prices in the table are the local-linear calls, floored, with puts by parity.

Bondarenko PCA. Start from the projected calls of Section 8.1. The density is the convolution of a Gaussian kernel on the spot with a discrete probability. The nodes are equally spaced from the first quoted strike to the last, and the spacing

is the bandwidth, with the count kept between 12 and 61. The weights are non-negative, sum to one, and have mean F . Each node prices a call by the Bachelier formula. The bandwidth minimises the even/odd out-of-the-money error. The selected widths are $h = 188.0, 364.1, 400.0$, and 171.3 . The NDX curve has four peaks. Prices in Table 7 are those calls, floored, with puts by parity.

Aït-Sahalia and Lo. Start from the projected calls of Section 8.1. On one maturity the futures price and the expiry are constant, and the semiparametric regression reduces to Nadaraya–Watson of Black–Scholes implied volatility on moneyness K/F . The kernel on that regressor is the order-2 Gaussian. **The bandwidth is not their Appendix A rule. Carried onto these chains that rule gives $h = 0.171, 0.303, 0.247$, and 0.315 , about one standard deviation of moneyness, and the smile is flat. The width used here is the smallest value on a geometric grid from 0.01 to 0.40 at which the density has one peak and total variation at most 1.05.** The selected widths are $h = 0.036, 0.025, 0.044$, and 0.036 . Prices are the Black–Scholes calls at the smoothed volatility, floored, with puts by parity. The density is the strike derivative of that call. Each curve has one peak and mass 1.0.

Grith et al. Start from the projected calls of Section 8.1. Both objects are a local cubic with the quartic kernel, and h is the half-width of that kernel. The bandwidth minimises the leave-one-out squared error of the local cubic, the cross-validation they compute and then set aside because the density is wiggly. The width in their comparison figure is twice that value: on the DAX chain, 114 against 228. The call row divides the call by the spot, reads the density as e^{rT} times the scaled second derivative, and prices from the local level. Cross-validation selects $h = 52.2, 317.9, 498.7$, and 110.7 , so the doubled widths are 104.4, 635.7, 997.4, and 221.5. The implied-volatility row is a local cubic of Black–Scholes implied volatility against the strike, with Black–Scholes prices at the local level and a density that uses the local slope and curvature. Its cross-validated widths are $h = 88.9, 504.8, 498.7$, and 204.0 , and the doubled widths are 177.9, 1009.6, 997.4, and 407.9. Their DAX figure for this row shows one widened curve, at 372, and does not print the cross-validated width, so the same doubling is used. The doubled call cubic is drawn on the left of Figure 3, with Aït-Sahalia and Lo. The cross-validated call cubic and both implied-volatility widths are wiggly. They are drawn in Appendix A so they do not crowd that figure. Prices are floored, with puts by parity.

Hold-out refits the projection and every tuned method on the training strikes

only, including the choice of its tuning constant, and scores the other half against the quoted mids. The two directions are averaged.

A fit inside the decreasing-convex cone cannot undercut the gap between the quoted calls and that cone. The plug-in local linear of Aït-Sahalia and Duarte remains far from the quotes. The positive convolution tracks them much more closely, and its NDX density has four peaks.

The estimator of this paper has the smallest hold-out in Table 7 on each chain, and the same ranking holds in the average. Each density has one peak, unit mass, and total variation at one. On three chains the selected penalty is zero; on NDX it is positive. In sample the estimator is below the positive convolution on each chain. The cross-validated cubics of Grith et al. are as tight or tighter in sample, and the hold-out reverses that ranking. The call cubic fails in the far wing, where a one-sided fit on the held-out strikes misses the put. Aït-Sahalia and Lo keeps one peak, with a larger pricing error. Once the tuning constants are known, that smile prices the full chain in the least time, and the plug-in local linear is faster than this estimator on each chain. The positive convolution is slower on each chain. The local polynomials take longer on the S&P 500 and Nasdaq chains. On the Russell chain they match this estimator to a tenth, and the cross-validated call cubic is slightly faster.

Total variation, defined in the caption of Table 7, charges an extra oscillation by its height relative to the mode. The cross-validated cubics oscillate. They are drawn in Appendix A rather than on the axes of Figure 3. The doubled call cubic, the plug-in local linear, and Aït-Sahalia and Lo are on the left of that figure.

Table 7: Out-of-the-money repricing. Every fit starts from the projection in Section 8.1 and is scored on the quoted mids. Ours is the smoothing spline (8), thriced at (15) ($h = 98.6, 165.9, 606.0, 59.0$). Ait-Sahalia and Duarte is the constrained local linear at the Fan–Gijbels plug-in; the mass is that density after the unit-mass scaling. **Ait-Sahalia and Lo does not use the Appendix A bandwidth**; the widths are the unimodal rule in that paragraph, $h = 0.036, 0.025, 0.044$, and 0.036 . Each Grith et al. object is shown at the leave-one-out width and at twice that width. The mass is the density of that same fit. PCA is priced from the Gaussian convolution on the spot. OTM is put RMSE for $K \leq F$ and call RMSE for $K \geq F$, premiums floored at zero. Mass is $\int \max(\hat{q}, 0)$ on $[\max(50, 0.2F), 2.4F]$. TV is the total variation of $\max(\hat{q}, 0)$ on $[0.55F, 1.40F]$, divided by twice the maximum; a unimodal curve scores 1.0. Hold-out refits the projection and the tuning on each half. Time is the wall-clock cost of a call at every quoted strike after the tuning constants are known, divided by Ours. Once the spline is chosen, those prices are read from one Fourier transform. Average is the equal-weight mean of the four chains. Ours is listed first. It has the smallest hold-out on each chain, with one peak and total variation 1.0. The other rows follow the SPX December hold-out, smallest first.

Chain	Method	Hold-out	RMSE OTM	Puts	Calls	Mass	TV	Time
SPX Dec	Ours	0.041	0.036	0.034	0.041	1.0	1.0	1.0
SPX Dec	PCA	0.055	0.051	0.053	0.048	1.0	1.0	13.8
SPX Dec	Grith et al., IV, 2×	0.055	0.032	0.031	0.037	1.0	1.1	4.6
SPX Dec	Grith et al., IV, CV	0.074	0.031	0.030	0.035	1.0	1.3	2.7
SPX Dec	Ait-Sahalia and Lo	4.626	4.626	3.442	7.124	1.0	1.0	0.4
SPX Dec	Ait-Sahalia and Duarte	5.385	4.272	3.263	6.450	1.0	1.0	0.6
SPX Dec	Grith et al., call, 2×	9.964	0.031	0.030	0.036	1.0	1.1	2.8
SPX Dec	Grith et al., call, CV	14.653	0.029	0.029	0.032	1.0	3.4	2.7
SPX Mar	Ours	0.066	0.056	0.052	0.067	1.0	1.0	1.0
SPX Mar	PCA	0.075	0.066	0.061	0.077	1.0	1.0	10.5
SPX Mar	Grith et al., IV, 2×	0.282	0.165	0.143	0.215	1.0	1.0	3.1
SPX Mar	Grith et al., IV, CV	0.091	0.051	0.049	0.057	1.0	1.0	2.9
SPX Mar	Ait-Sahalia and Lo	1.312	0.910	0.609	1.474	1.0	1.1	0.3
SPX Mar	Ait-Sahalia and Duarte	9.965	7.837	5.633	12.202	1.0	1.0	0.3
SPX Mar	Grith et al., call, 2×	1.339	0.152	0.107	0.240	1.0	1.0	2.9
SPX Mar	Grith et al., call, CV	18.635	0.042	0.040	0.048	1.0	1.0	2.7
NDX Dec	Ours	1.109	1.065	1.113	0.792	1.0	1.0	1.0
NDX Dec	PCA	1.324	1.151	1.117	1.304	1.0	1.3	19.1
NDX Dec	Grith et al., IV, 2×	1.236	0.946	0.991	0.686	1.0	1.2	2.5
NDX Dec	Grith et al., IV, CV	1.248	0.867	0.921	0.524	1.0	3.0	1.7
NDX Dec	Ait-Sahalia and Lo	25.840	25.839	21.836	40.004	1.0	1.0	0.2
NDX Dec	Ait-Sahalia and Duarte	19.189	16.023	14.909	20.619	1.0	1.1	0.3
NDX Dec	Grith et al., call, 2×	40.791	0.946	0.993	0.677	1.0	1.1	1.7
NDX Dec	Grith et al., call, CV	79.411	0.867	0.922	0.523	0.9	2.8	1.7
RUT Dec	Ours	0.031	0.026	0.027	0.024	1.0	1.0	1.0
RUT Dec	PCA	0.067	0.058	0.066	0.043	1.0	1.0	3.6
RUT Dec	Grith et al., IV, 2×	0.116	0.042	0.047	0.033	1.3	1.0	1.0
RUT Dec	Grith et al., IV, CV	0.039	0.026	0.027	0.025	1.0	1.1	1.0
RUT Dec	Ait-Sahalia and Lo	1.117	1.117	1.262	0.870	1.0	1.0	0.1
RUT Dec	Ait-Sahalia and Duarte	2.834	2.224	1.968	2.544	1.0	1.1	0.1
RUT Dec	Grith et al., call, 2×	0.432	0.056	0.051	0.063	1.0	1.0	1.0
RUT Dec	Grith et al., call, CV	7.026	0.022	0.022	0.023	1.0	1.1	0.9
Average	Ours	0.312	0.296	0.306	0.231	1.0	1.0	1.0
Average	PCA	0.380	0.331	0.324	0.368	1.0	1.1	11.7
Average	Grith et al., IV, 2×	0.422	0.296	0.303	0.243	1.1	1.1	2.8
Average	Grith et al., IV, CV	0.363	0.244	0.257	0.160	1.0	1.6	2.1
Average	Ait-Sahalia and Lo	8.224	8.123	6.787	12.368	1.0	1.0	0.2
Average	Ait-Sahalia and Duarte	9.343	7.589	6.443	10.454	1.0	1.0	0.3
Average	Grith et al., call, 2×	13.131	0.297	0.295	0.254	1.0	1.1	2.1
Average	Grith et al., call, CV	29.932	0.240	0.253	0.156	1.0	2.1	2.0




Figure 3: SPX, NDX, and RUT, 18 December 2026, snapshot 23 September 2026. Each curve is fit to the projected call of Section 8.1. Left: $\max(\hat{f}_Q, 0)$ for the smoothing spline at (15), the Gaussian convolution of Bondarenko, the call cubic of Grith et al. at twice its cross-validated bandwidth, the constrained local linear of Ait-Sahalia and Duarte at the Fan–Gijbels plug-in, and Ait-Sahalia and Lo at the bandwidth in Section 8.2. Right: out-of-the-money implied volatility for the same five fits. The dashed line is the forward. The cross-validated call cubic and the implied-volatility cubics are wiggly, and are drawn in Appendix A so they do not crowd this figure.

9. Conclusion

The estimator reads the risk-neutral density off the empirical measure of a single maturity. Calls are rescaled to a complementary distribution, a smoothing spline fits that measure, the unquoted tails are filled, and a thriced Gaussian convolution of the spline is the density and the call. The penalty is zero when the thriced density is already unimodal. On exact quotes that is the selected value, so those

rows are the thrice-d spline at (15). The estimator is the smallest on sparse Heston and on dense variance gamma. On dense Heston the cross-validated implied-volatility cubic and the Nadaraya–Watson smile are smaller, and on sparse variance gamma that smile is smaller. With noise in implied volatility that smile is smaller on the dense chain, and on the sparse chain the implied-volatility cubics and that smile are smaller. On the four listed chains the hold-out error is smaller than every other method on each chain, and each density has one peak, unit mass, and total variation at one. The cross-validated cubics are as tight or tighter in sample. The call cubic does not hold out on the far wing.

Appendix A. Other listed densities

The cross-validated call cubic of Grith et al., and that paper’s implied-volatility cubic at the leave-one-out width and at twice that width, oscillate inside the window of Figure 3. Figure A.4 draws them on the same three chains. They are kept off the axes of Figure 3 so that figure is not crowded with those wiggles. Aït-Sahalia and Lo is on Figure 3.




Figure A.4: Other kernels on the chains of Figure 3. The call cubic of Grith et al. is at the leave-one-out width, and the implied-volatility cubic is drawn at that width and at twice that width. Left: $\max(\hat{f}_Q, 0)$, on the strike window and vertical scale of the left panel of Figure 3, with curves truncated there. Right: out-of-the-money implied volatility of the same fits. The dashed line is the forward.

Code and Data Availability

A Python implementation of (14) and a replication script are available at <https://github.com/s-broda/RND>.

References

Aït-Sahalia, Y., Duarte, J., 2003. Nonparametric option pricing under shape restrictions. *Journal of Econometrics* 116, 9–47.

- Aït-Sahalia, Y., Lo, A.W., 1998. Nonparametric estimation of state-price densities implicit in financial asset prices. *Journal of Finance* 53, 499–547.
- Aït-Sahalia, Y., Lo, A.W., 2000. Nonparametric risk management and implied risk aversion. *Journal of Econometrics* 94, 9–51.
- Albrecher, H., Mayer, P., Schoutens, W., Tistaert, J., 2007. The little Heston trap. *Wilmott Magazine*, 83–92.
- Anagnou, I., Bedendo, M., Hodges, S., Tompkins, R., 2005. Forecasting accuracy of implied and GARCH-based probability density functions. *Review of Futures Markets* 11, 41–66.
- Bahra, B., 1997. Implied Risk-Neutral Probability Density Functions from Option Prices: Theory and Application. Technical Report Working Paper 66. Bank of England.
- Bakshi, G., Cao, C., Chen, Z., 1997. Empirical performance of alternative option pricing models. *Journal of Finance* 52, 2003–2049.
- Birke, M., Pilz, K.F., 2009. Nonparametric option pricing with no-arbitrage constraints. *Journal of Financial Econometrics* 7, 53–76.
- Bliss, R.R., Panigirtzoglou, N., 2002. Testing the stability of implied probability density functions. *Journal of Banking & Finance* 26, 381–422.
- Bondarenko, O., 2003. Estimation of risk-neutral densities using positive convolution approximation. *Journal of Econometrics* 116, 85–112.
- Breeden, D.T., Litzenberger, R.H., 1978. Prices of state-contingent claims implicit in option prices. *Journal of Business* 51, 621–651.
- Buchen, P.W., Kelly, M., 1996. The maximum entropy distribution of an asset inferred from option prices. *Journal of Financial and Quantitative Analysis* 31, 143–159.
- Carr, P., Madan, D., 1999. Option valuation using the fast Fourier transform. *Journal of Computational Finance* 2, 61–73.
- Cont, R., Tankov, P., 2004. Non-parametric calibration of jump-diffusion option pricing models. *Journal of Computational Finance* 7, 1–49.

- Cooper, N., 1999. Testing techniques for estimating implied RNDs from the prices of European-style options, in: *Proceedings of the BIS Workshop on Estimating and Interpreting Probability Density Functions*. Bank for International Settlements.
- Corrado, C.J., Su, T., 1996. Skewness and kurtosis in S&P 500 index returns implied by option prices. *Journal of Financial Research* 19, 175–192.
- Coutant, S., Jondeau, E., Rockinger, M., 2001. Reading PIBOR futures options smiles: The 1997 snap election. *Journal of Banking & Finance* 25, 1957–1987.
- Dupire, B., 1994. Pricing with a smile. *Risk* 7, 18–20.
- Fan, J., Gijbels, I., 1996. *Local Polynomial Modelling and Its Applications*. Chapman & Hall, London.
- Fan, J., Mancini, L., 2009. Option pricing with model-guided nonparametric methods. *Journal of the American Statistical Association* 104, 1351–1372.
- Fengler, M.R., 2005. *Semiparametric Modeling of Implied Volatility*. Springer, Berlin.
- Fengler, M.R., 2009. Arbitrage-free smoothing of the implied volatility surface. *Quantitative Finance* 9, 417–428.
- Figlewski, S., 2010. Estimating the implied risk neutral density for the U.S. market portfolio, in: Bollerslev, T., Russell, J., Watson, M. (Eds.), *Volatility and Time Series Econometrics: Essays in Honor of Robert Engle*. Oxford University Press, pp. 323–358.
- Figlewski, S., 2018. Risk-neutral densities: A review. *Annual Review of Financial Economics* 10, 329–359.
- Gasser, T., Müller, H.G., 1979. Kernel estimation of regression functions, in: Gasser, T., Rosenblatt, M. (Eds.), *Smoothing Techniques for Curve Estimation*. Springer, Berlin, pp. 23–68.
- Gatheral, J., Jacquier, A., 2014. Arbitrage-free SVI volatility surfaces. *Quantitative Finance* 14, 59–71.

- Grith, M., Härdle, W., Schienle, M., 2012. Nonparametric estimation of risk-neutral densities, in: Duan, J.C., Härdle, W.K., Gentle, J.E. (Eds.), *Handbook of Computational Finance*. Springer, pp. 277–305.
- Heston, S.L., 1993. A closed-form solution for options with stochastic volatility with applications to bond and currency options. *Review of Financial Studies* 6, 327–343.
- Jackwerth, J.C., 2004. Option-Implied Risk-Neutral Distributions and Risk Aversion. Research Foundation of AIMR, Charlottesville.
- Jackwerth, J.C., Rubinstein, M., 1996. Recovering probability distributions from option prices. *Journal of Finance* 51, 1611–1631.
- Jarrow, R., Rudd, A., 1982. Approximate option valuation for arbitrary stochastic processes. *Journal of Financial Economics* 10, 347–369.
- Jondeau, E., Rockinger, M., 2000. Reading the smile: The message conveyed by methods which infer risk neutral densities. *Journal of International Money and Finance* 19, 885–915.
- Lai, W.N., 2014. Comparison of methods to estimate option implied risk-neutral densities. *Quantitative Finance* 14, 1839–1855.
- Lewis, A.L., 2000. *Option Valuation under Stochastic Volatility*. Finance Press, Newport Beach.
- Liu, X., Shackleton, M.B., Taylor, S.J., Xu, X., 2007. Closed-form transformations from risk-neutral to real-world distributions. *Journal of Banking & Finance* 31, 1501–1520.
- Lu, S., 2019. Monte Carlo analysis of methods for extracting risk-neutral densities with affine jump diffusions. *Journal of Futures Markets* 39, 1587–1612.
- Ludwig, M., 2015. Robust estimation of shape-constrained state price density surfaces. *Journal of Derivatives* 22, 56–72.
- Madan, D.B., Carr, P.P., Chang, E.C., 1998. The variance gamma process and option pricing. *European Finance Review* 2, 79–105.
- Malz, A.M., 1997. Estimating the probability distribution of the future exchange rate from option prices. *Journal of Derivatives* 5, 18–36.

- Melick, W.R., Thomas, C.P., 1997. Recovering an asset's implied PDF from option prices: An application to crude oil during the Gulf crisis. *Journal of Financial and Quantitative Analysis* 32, 91–115.
- Monnier, J.B., 2014. Risk-neutral density recovery via spectral analysis. *SIAM Journal on Financial Mathematics* 5, 558–587. doi:10.1137/110840340.
- Ritchey, R.J., 1990. Call option valuation for discrete normal mixtures. *Journal of Financial Research* 13, 285–296.
- Rompolis, L.S., Tzavalis, E., 2007. Retrieving risk neutral densities based on risk neutral moments through a Gram–Charlier series expansion. *Mathematical and Computer Modelling* 46, 225–234.
- Ross, S.A., 1976. Options and efficiency. *Quarterly Journal of Economics* 90, 75–89.
- Rubinstein, M., 1994. Implied binomial trees. *Journal of Finance* 49, 771–818.
- Santos, A.A.P., Guerra, J., 2015. Implied risk neutral densities from option prices: Hypergeometric, spline, lognormal, and Edgeworth functions. *Journal of Futures Markets* 35, 655–678.
- Schucany, W.R., Sommers, J.P., 1977. Improvement of kernel type density estimators. *Journal of the American Statistical Association* 72, 420–423.
- Shimko, D., 1993. Bounds of probability. *Risk* 6, 33–37.
- Söderlind, P., Svensson, L.E.O., 1997. New techniques to extract market expectations from financial instruments. *Journal of Monetary Economics* 40, 383–429.
- Stutzer, M., 1996. A simple nonparametric approach to derivative security valuation. *Journal of Finance* 51, 1633–1652.
- Xiu, D., 2014. Hermite polynomial based expansion of European option prices. *Journal of Econometrics* 179, 158–177.
- Yatchew, A., Härdle, W., 2006. Nonparametric state price density estimation using constrained least squares and the bootstrap. *Journal of Econometrics* 133, 579–609.